

Practical 1 and 2 - Retake

Melvil Deleage & Jeff Macaraeg

June 15, 2025

Table of contents

Introduction	3
Practical 1, Part 1 - Financial Returns & Normality	4
Question a)	4
Question b)	4
Question c)	5
Question d)	6
Question e)	7
Question f)	8
Question g)	9
Practical 1, Part 2 - Financial time series, volatility and the random walk hypothesis	10
Question a)	10
Question b)	10
Question c)	11
Question d)	13
Question e)	14
Question f)	15
Question g)	21
Question h)	24
Practical 2, Part 1 - Venice	29
Question a)	29
Question b)	30
Question c)	31
Question d)	32
Question e)	33
Question f) + g)	33
Question h)	34
Question i)	34
Practical 2, Part 2 - Nuclear Reactors	34
Question a)	34
Question b)	35
Question c)	36
Question d)	36
Question e)	37
Question f)	37

Practical 2, Part 3 - Night temperatures in Lausanne	38
Question a)	38
Question b)	38
Question c)	39
Question d)	39
Question e)	40

Introduction

This documents is the annex to our final report for the retake project. It contains the full code for the practicals 1 and 2 of the retake with the results and some discussions. This allows anyone to reproduce our results and to understand the code we used.

Practical 1, Part 1 - Financial Returns & Normality

Question a)

Read in the data. Then, assess the stationarity of the (raw) stock indices.

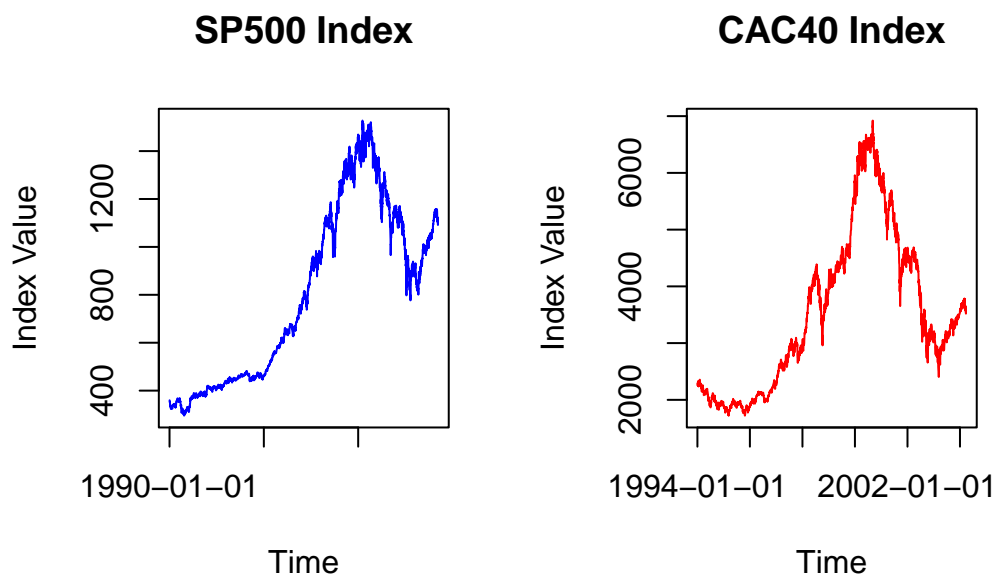
As discussed with the teacher, we will only select 2 indices: SP500 & CAC40.

Augmented Dickey-Fuller Test

```
data: sp500_ts
Dickey-Fuller = -1.2128, Lag order = 15, p-value = 0.9044
alternative hypothesis: stationary
```

Augmented Dickey-Fuller Test

```
data: cac40_ts
Dickey-Fuller = -0.7332, Lag order = 13, p-value = 0.9676
alternative hypothesis: stationary
```




As a result, we can see that for SP500, the p-value is 0.9044 and for CAC40, the p-value is 0.9676. As they are both higher than 0.05, we can conclude that both series are not stationary.

Question b)

Create a function to transform the daily stock indices into their daily negative log returns counterparts. Plot the latter series and assess their stationarity. To compare the series, also plot the negative log returns on a common scale to all indices.

SP500 Negative Log Return CAC40 Negative Log Return



```
[1] "ADF Test for SP500:"
```

```
Augmented Dickey-Fuller Test
```

```
data:  sp500_ret
Dickey-Fuller = -14.843, Lag order = 15, p-value = 0.01
alternative hypothesis: stationary
```

```
[1] "ADF Test for CAC40:"
```

```
Augmented Dickey-Fuller Test
```

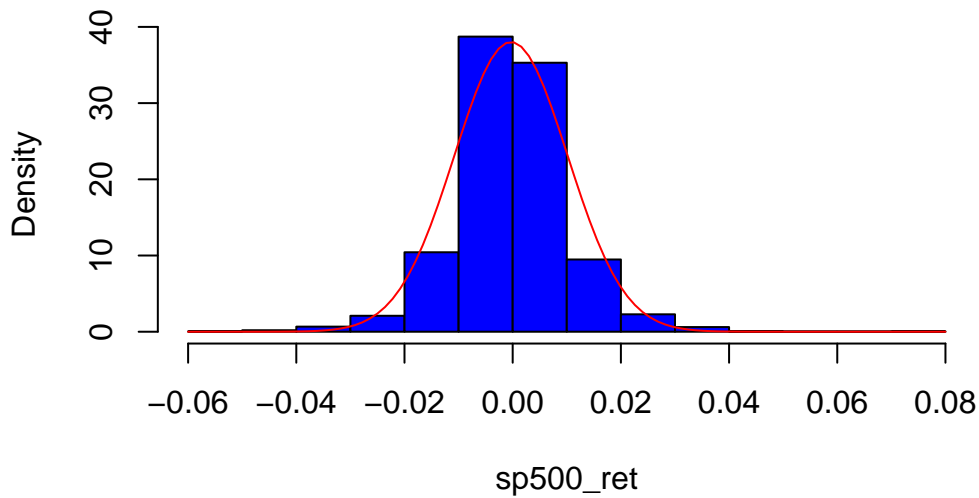
```
data:  cac40_ret
Dickey-Fuller = -13.34, Lag order = 13, p-value = 0.01
alternative hypothesis: stationary
```

As both plots seem to not show trends, periodic cycles and a stable variance, the series seem to be stationary. To verify this, the ADF test shows that both p-values (p-value = 0.01) are lower than 0.05, so we reject the hypothesis and thus, both series are stationary.

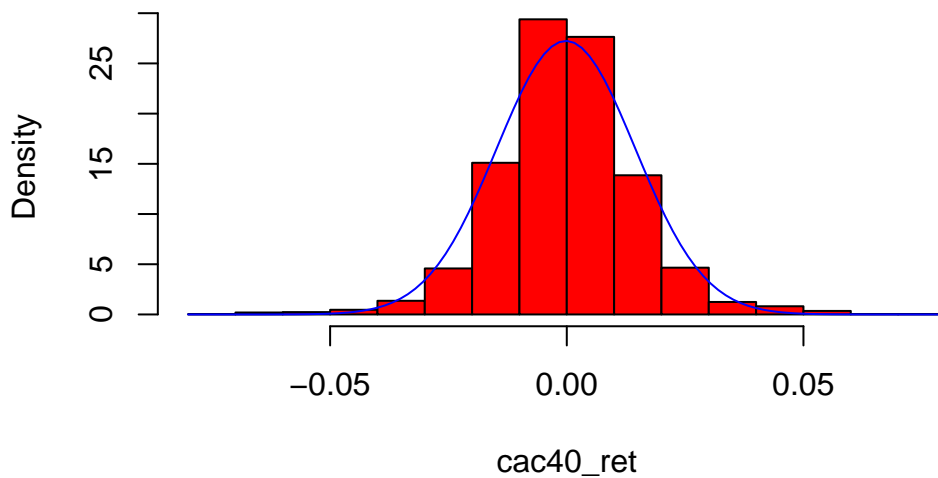
Question c)

Draw histograms of the negative log returns and compare them to the Normal distribution. What do you observe?

Histogram of SP500 Returns



Histogram of CAC40 Returns




Both histograms have bell-shaped distributions, but are not perfectly aligned with the normal curve. Also, the tails seem to go further than what the normal distribution curve predicts, which can indicate a higher probability of extreme returns than expected in a normal model.


Question d)

Check the normality assumption of the negative log returns using QQ-plots. What is your conclusion?

QQ-plot SP500 Returns



QQ-plot CAC40 Returns



On both the lower and upper extremes, the points deviate from the line, showing heavier tails than a normal distribution. So, the negative log returns don't seem to follow normal distribution.

Question e)

Formally test the normality assumption of the negative log returns using an Anderson-Darling testing procedure. Do you reject the Normal hypothesis?

Anderson-Darling normality test

```
data: sp500_ret  
A = 29.24, p-value < 2.2e-16
```

Anderson-Darling normality test

```
data: cac40_ret
A = 10.33, p-value < 2.2e-16
```

As both p-values are lower than 0.05 (both at 2.2e-16), we reject the null hypothesis, meaning that the returns are not normally distributed.

Question f)

Use the `fitdistr()` function from the MASS package in order to obtain the (maximum-likelihood estimated) parameters of distributions you could imagine for the negative log returns. Try to fit at least two different distributions on the data and, using an information criteria (such as the AIC), decide which distributional framework fits best for each of the series.

```
      mean      sd
-0.0003139511  0.0104904994
( 0.0001751582) ( 0.0001238556)
```

```
      m      s      df
-0.0003823438  0.0083396816  6.5020438529
( 0.0001642079) ( 0.0011490812) ( 5.6443703001)
```

```
      mean      sd
-0.0001723074  0.0146403328
( 0.0002884550) ( 0.0002039685)
```

```
      m      s      df
-0.0003374078  0.0116334321  5.2499036218
( 0.0002632892) ( 0.0002745656) ( 0.5356059623)
```

```
[1] -22510.5
```

```
[1] -22898.7
```

```
[1] -14447.55
```

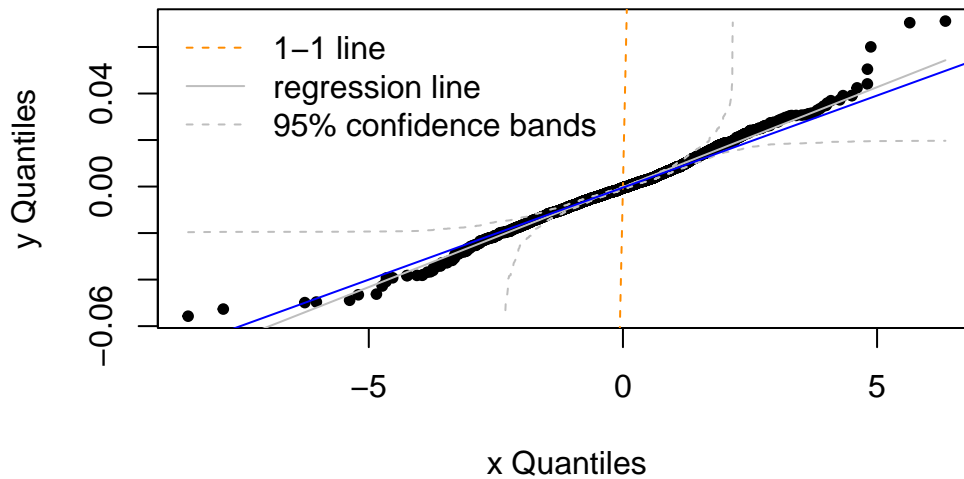
```
[1] -14626
```

As the Student model yields lower AIC than the normal model, it indicates a better balance between the goodness-of-fit and the complexity. It is a more appropriate choice of model.

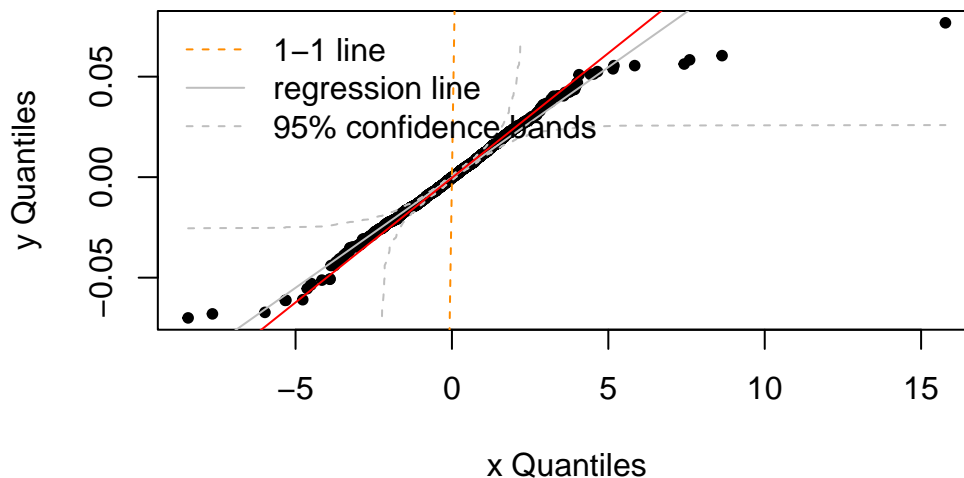
Question g)

If this has not been done in (f), fit a t-distribution to the negative log returns using `fitdistr()`. Using a QQ-plot for each of the series, decide whether the fit is better than with a Normal distribution, based on your answer in (d).

QQ-plot SP500 vs t-Distribution



QQ-plot CAC40 vs t-Distribution

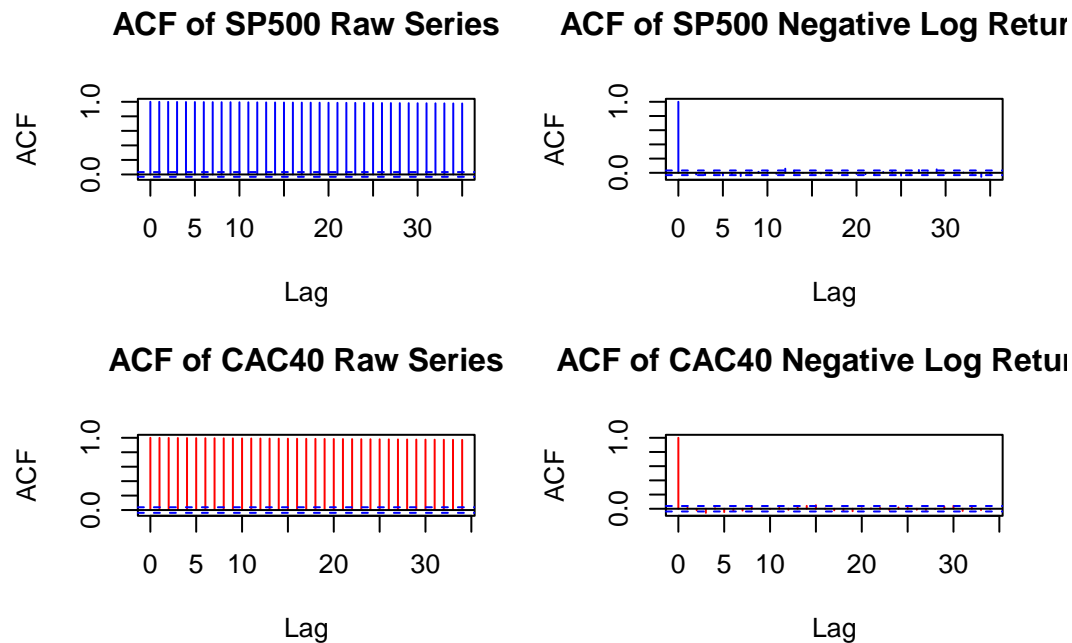


Despite some deviations from points in the extremes, the QQ-plots confirm that a Student's t-distribution provides a better fit for both the SP500 and CAC40 log returns compared to a normal distribution. The heavier tails of the t-distribution captures the extremes better. It is also consistent with question d).

Practical 1, Part 2 - Financial time series, volatility and the random walk hypothesis

Question a)

Plot the ACF of all the series in Part 1 (i.e. the raw series as well as the negative log returns). What do you observe?



The raw series show high autocorrelation at all lags, implying non-stationarity. For the negative log returns, it shows little to no autocorrelation, suggesting that the daily returns are close to uncorrelated.

Question b)

Use a Ljung-Box procedure to formally test for (temporal) serial dependence in the series. What is your conclusion?

Box-Ljung test

```
data: sp500
X-squared = 71035, df = 20, p-value < 2.2e-16
```

Box-Ljung test

```
data: sp500_ret
X-squared = 38.931, df = 20, p-value = 0.0068
```

Box-Ljung test

```
data: cac40
X-squared = 50889, df = 20, p-value < 2.2e-16
```

Box-Ljung test


```
data: cac40_ret  
X-squared = 41.079, df = 20, p-value = 0.003639
```

Looking at the raw series for both CAC40 and SP500, we obtain a p-value that is lower than 0.05 (both at $2.2e-16$), so we reject the null hypothesis and the raw series show autocorrelation. For both negative log returns, the p-values are higher than the raw series (0.0068 for SP500, 0.003639 for CAC40), but still smaller than the p-value, so again they show autocorrelation.


Question c)

Propose ARIMA models for each of the negative log returns series, based on visualisation tools (e.g. ACF and PACF). Select an ARIMA model using `auto.arima()` (forecast package) to each of the negative log returns series. Comment on the difference. Assess the residuals of the resulting models.

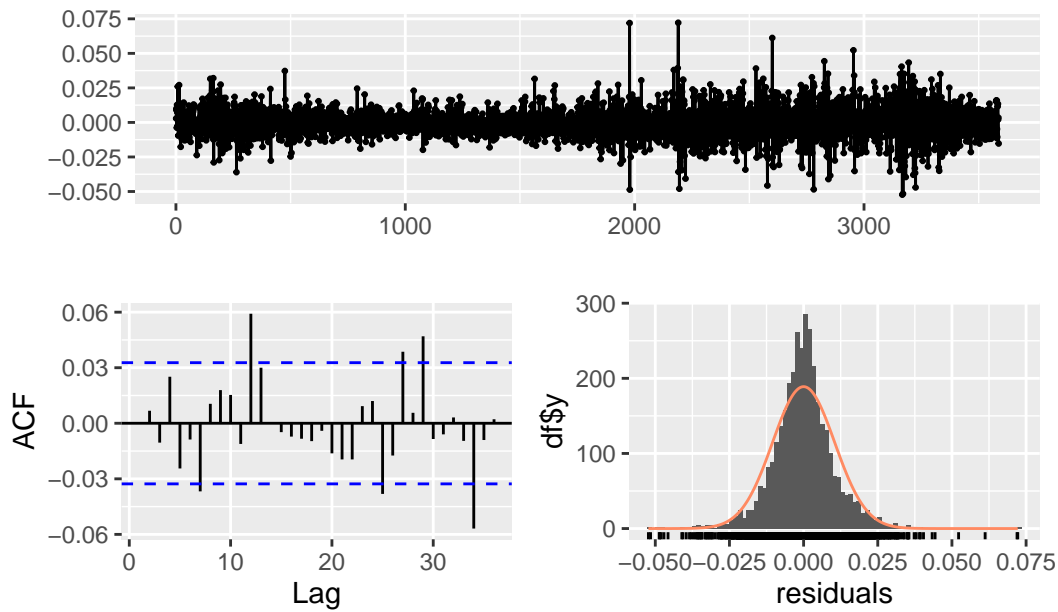
ACF – SP500 Negative Log Return PACF – SP500 Negative Log Return



ACF – CAC40 Negative Log Return PACF – CAC40 Negative Log Return



Residuals from ARIMA(2,0,1) with non-zero mean




Ljung-Box test

data: Residuals from ARIMA(2,0,1) with non-zero mean
 $Q^* = 12.498$, $df = 7$, $p\text{-value} = 0.08534$

Model df: 3. Total lags used: 10

Residuals from ARIMA(0,0,0) with zero mean



Ljung-Box test

data: Residuals from ARIMA(0,0,0) with zero mean
 $Q^* = 21.191$, $df = 10$, $p\text{-value} = 0.0198$

Model df: 0. Total lags used: 10

Series: sp500_ret
ARIMA(2,0,1) with non-zero mean

Coefficients:

	ar1	ar2	ma1	mean
	0.7796	-0.0291	-0.7827	-3e-04
s.e.	0.1035	0.0180	0.1024	2e-04

$\sigma^2 = 0.0001099$: log likelihood = 11261.47
AIC=-22512.94 AICc=-22512.92 BIC=-22482.01

Training set error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	2.252009e-06	0.01047816	0.00748605	-Inf	Inf	0.6910122

ACF1

Training set -0.0003543668

Series: cac40_ret
ARIMA(0,0,0) with zero mean

$\sigma^2 = 0.0002144$: log likelihood = 7225.6
AIC=-14449.19 AICc=-14449.19 BIC=-14443.34

Training set error measures:


	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Training set	-0.0001723074	0.01464135	0.01089581	100	100	0.7046982	0.01343285

The ACF & PACF for both SP500 & CAC40 negative log returns show almost no significant autocorrelation, indicating that the returns behave in a similar way to white noise. It is also confirmed by checking the residuals. We can also observe that the `auto.arima()` function selected very simple models.

Question d)

Assess the residuals of the resulting models from (c), both their raw values and their absolute values, through visual tools (such as the ACF) and formal tests (e.g. Ljung-Box). What do you conclude about the independence assumption?


Series residuals(arima_sp500) Series residuals(arima_cac40)




The independence assumption holds for the raw residuals but not for the volatility patterns (autocorrelation in the absolute residuals). So, further modeling could be used to fully capture the dynamics of the residuals.

Question e)

Plot the volatility of the raw series of indices. What is your conclusion on the homoscedasticity assumption?



CAC40 Volatility



The data shows heteroskedasticity, as the volatility (or variance) is not constant over time.

Question f)

Fit GARCH models to the negative log returns of each series with both standardised and skewed t-distributions, with order (1, 1), using the `garchFit()` function from the `fGarch` library. Assess the quality of the fit by evaluating the residuals.

Title:

GARCH Modelling

Call:

```
garchFit(formula = ~garch(1, 1), data = sp500_ret, cond.dist = "std",
  trace = FALSE)
```

Mean and Variance Equation:

```
data ~ garch(1, 1)
<environment: 0x130b340c>
[data = sp500_ret]
```

Conditional Distribution:

std

Coefficient(s):

mu	omega	alpha1	beta1	shape
-5.6197e-04	3.2345e-07	5.0937e-02	9.4765e-01	7.1692e+00

Std. Errors:

based on Hessian

Error Analysis:

Estimate	Std. Error	t value	Pr(> t)
----------	------------	---------	----------

```

mu      -5.620e-04  1.250e-04  -4.497  6.90e-06 ***
omega   3.234e-07  1.442e-07   2.242   0.0249 *
alpha1  5.094e-02  7.720e-03   6.598  4.16e-11 ***
beta1   9.476e-01  7.660e-03  123.714 < 2e-16 ***
shape   7.169e+00  8.189e-01   8.755 < 2e-16 ***
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Log Likelihood:

11798.72 normalized: 3.289301

Description:

Sun Jun 15 09:52:16 2025 by user:

Standardised Residuals Tests:

			Statistic	p-Value
Jarque-Bera Test	R	Chi^2	661.3620424	0.000000000
Shapiro-Wilk Test	R	W	0.9842328	0.000000000
Ljung-Box Test	R	Q(10)	18.0504009	0.054119378
Ljung-Box Test	R	Q(15)	34.5937376	0.002808328
Ljung-Box Test	R	Q(20)	36.2946867	0.014198662
Ljung-Box Test	R^2	Q(10)	7.7661792	0.651664153
Ljung-Box Test	R^2	Q(15)	10.7404577	0.770766521
Ljung-Box Test	R^2	Q(20)	15.4320193	0.751176895
LM Arch Test	R	TR^2	10.0402112	0.612432829

Information Criterion Statistics:

AIC	BIC	SIC	HQIC
-6.575815	-6.567193	-6.575819	-6.572742

Title:

GARCH Modelling

Call:

```
garchFit(formula = ~garch(1, 1), data = sp500_ret, cond.dist = "sstd",
  trace = FALSE)
```

Mean and Variance Equation:

data ~ garch(1, 1)

<environment: 0x122f93e20>

[data = sp500_ret]

Conditional Distribution:

sstd

Coefficient(s):

mu	omega	alpha1	beta1	skew	shape
-5.1103e-04	3.3855e-07	5.1496e-02	9.4679e-01	1.0316e+00	7.3173e+00

Std. Errors:

based on Hessian

Error Analysis:

	Estimate	Std. Error	t value	Pr(> t)
mu	-5.110e-04	1.308e-04	-3.906	9.4e-05 ***
omega	3.386e-07	1.472e-07	2.299	0.0215 *
alpha1	5.150e-02	7.770e-03	6.628	3.4e-11 ***
beta1	9.468e-01	7.772e-03	121.816	< 2e-16 ***
skew	1.032e+00	2.406e-02	42.875	< 2e-16 ***
shape	7.317e+00	8.596e-01	8.512	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Log Likelihood:

11799.62 normalized: 3.28955

Description:

Sun Jun 15 09:52:17 2025 by user:

Standardised Residuals Tests:

			Statistic	p-Value
Jarque-Bera Test	R	Chi ²	659.5154771	0.000000000
Shapiro-Wilk Test	R	W	0.9842608	0.000000000
Ljung-Box Test	R	Q(10)	17.9330634	0.056103319
Ljung-Box Test	R	Q(15)	34.5128308	0.002883347
Ljung-Box Test	R	Q(20)	36.1922613	0.014599763
Ljung-Box Test	R ²	Q(10)	7.5116696	0.676416753
Ljung-Box Test	R ²	Q(15)	10.5335043	0.784908436
Ljung-Box Test	R ²	Q(20)	15.1615192	0.767092017
LM Arch Test	R	TR ²	9.8488270	0.629221076

Information Criterion Statistics:

AIC	BIC	SIC	HQIC
-6.575755	-6.565409	-6.575761	-6.572067

Title:

GARCH Modelling

Call:

```
garchFit(formula = ~garch(1, 1), data = cac40_ret, cond.dist = "std",
  trace = FALSE)
```

Mean and Variance Equation:

```
data ~ garch(1, 1)
<environment: 0x121c39318>
[data = cac40_ret]
```

Conditional Distribution:

std

Coefficient(s):

mu	omega	alpha1	beta1	shape
----	-------	--------	-------	-------

-5.2677e-04 2.0876e-06 6.4925e-02 9.2750e-01 1.0000e+01

Std. Errors:

based on Hessian

Error Analysis:

	Estimate	Std. Error	t value	Pr(> t)
mu	-5.268e-04	2.343e-04	-2.248	0.02459 *
omega	2.088e-06	7.696e-07	2.712	0.00668 **
alpha1	6.493e-02	9.599e-03	6.764	1.35e-11 ***
beta1	9.275e-01	1.043e-02	88.943	< 2e-16 ***
shape	1.000e+01	1.361e+00	7.350	1.99e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Log Likelihood:

7496.61 normalized: 2.910175

Description:

Sun Jun 15 09:52:16 2025 by user:

Standardised Residuals Tests:

			Statistic	p-Value
Jarque-Bera Test	R	Chi^2	36.1606492	1.405448e-08
Shapiro-Wilk Test	R	W	0.9968239	3.245060e-05
Ljung-Box Test	R	Q(10)	12.1818800	2.730684e-01
Ljung-Box Test	R	Q(15)	20.7654478	1.444922e-01
Ljung-Box Test	R	Q(20)	22.3308637	3.228303e-01
Ljung-Box Test	R^2	Q(10)	12.6755294	2.423833e-01
Ljung-Box Test	R^2	Q(15)	13.1950834	5.872327e-01
Ljung-Box Test	R^2	Q(20)	14.8593966	7.843956e-01
LM Arch Test	R	TR^2	13.7118254	3.194879e-01

Information Criterion Statistics:

AIC	BIC	SIC	HQIC
-5.816467	-5.805105	-5.816475	-5.812348

Title:

GARCH Modelling

Call:

```
garchFit(formula = ~garch(1, 1), data = cac40_ret, cond.dist = "sstd",
  trace = FALSE)
```

Mean and Variance Equation:

data ~ garch(1, 1)

<environment: 0x121f44390>

[data = cac40_ret]

Conditional Distribution:

sstd

Coefficient(s):

mu	omega	alpha1	beta1	skew	shape
-4.1284e-04	2.0402e-06	6.5401e-02	9.2730e-01	1.0840e+00	1.0000e+01

Std. Errors:

based on Hessian

Error Analysis:

	Estimate	Std. Error	t value	Pr(> t)
mu	-4.128e-04	2.377e-04	-1.737	0.08245 .
omega	2.040e-06	7.576e-07	2.693	0.00708 **
alpha1	6.540e-02	9.543e-03	6.853	7.23e-12 ***
beta1	9.273e-01	1.030e-02	90.004	< 2e-16 ***
skew	1.084e+00	3.327e-02	32.581	< 2e-16 ***
shape	1.000e+01	1.342e+00	7.452	9.24e-14 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Log Likelihood:

7500.141 normalized: 2.911545

Description:

Sun Jun 15 09:52:18 2025 by user:

Standardised Residuals Tests:


			Statistic	p-Value
Jarque-Bera Test	R	Chi^2	36.6432206	1.104144e-08
Shapiro-Wilk Test	R	W	0.9967943	2.934125e-05
Ljung-Box Test	R	Q(10)	12.0943005	2.787949e-01
Ljung-Box Test	R	Q(15)	20.7596486	1.446860e-01
Ljung-Box Test	R	Q(20)	22.3409037	3.223024e-01
Ljung-Box Test	R^2	Q(10)	12.3893095	2.598459e-01
Ljung-Box Test	R^2	Q(15)	12.9276394	6.078871e-01
Ljung-Box Test	R^2	Q(20)	14.5756952	8.001500e-01
LM Arch Test	R	TR^2	13.4251438	3.389111e-01

Information Criterion Statistics:


AIC	BIC	SIC	HQIC
-5.818432	-5.804797	-5.818443	-5.813490

residuals(garch_sp500_t)

Residuals: SP500 – t




ACF Residuals: SP500 – t




residuals(garch_sp500_skt)

Residuals: SP500 – skewed t




ACF Residuals: SP500 – skewed




residuals(garch_cac40_t)

Residuals: CAC40 – t




ACF Residuals: CAC40 – t




residuals(garch_cac40_skt)

Residuals: CAC40 – skewed t



ACF Residuals: CAC40 – skewec



Box-Ljung test

```
data: residuals(garch_sp500_t)
X-squared = 38.931, df = 20, p-value = 0.0068
```

Box-Ljung test

```
data: residuals(garch_sp500_skt)
X-squared = 38.931, df = 20, p-value = 0.0068
```

Box-Ljung test

```
data: residuals(garch_cac40_t)
X-squared = 41.079, df = 20, p-value = 0.003639
```

Box-Ljung test

```
data: residuals(garch_cac40_skt)
X-squared = 41.079, df = 20, p-value = 0.003639
```

By evaluating the residuals, we observe that for both SP500 and CAC40, the p-values are below 0.05, meaning there is still autocorrelation and residuals are not white noise. Looking at the AIC/BIC and the log-likelihood, it seems that the skewed distribution is slightly better for both models especially for CAC40. Thus, the Garch models are a good start but can be improved.

Question g)

Residual serial correlation can be present when fitting a GARCH directly on the negative log returns. Hence, in order to circumvent this problem, it is possible to use the following two-step approach:

- fit an ARMA(p,q) on the negative log returns;
- fit a GARCH(1,1) on the residuals of the ARMA(p,q) fit. Proceed with the above recipe. Assess the quality of the above fit.

To fit an ARMA(p,q) on the negative log returns:

To fit GARCH(1,1) on the ARMA residuals:

To assess the fit quality:

Title:

GARCH Modelling

Call:

```
garchFit(formula = ~garch(1, 1), data = res_sp500, cond.dist = "sstd",
  trace = FALSE)
```

Mean and Variance Equation:

```
data ~ garch(1, 1)
<environment: 0x134c54cf8>
[data = res_sp500]
```

Conditional Distribution:

sstd

Coefficient(s):

mu	omega	alpha1	beta1	skew	shape
-2.2520e-05	3.3032e-07	5.0301e-02	9.4808e-01	1.0551e+00	7.1935e+00

Std. Errors:

based on Hessian

Error Analysis:

	Estimate	Std. Error	t value	Pr(> t)
mu	-2.252e-05	1.325e-04	-0.170	0.8651
omega	3.303e-07	1.438e-07	2.298	0.0216 *

```
alpha1  5.030e-02   7.540e-03    6.671 2.53e-11 ***
beta1   9.481e-01   7.530e-03  125.907 < 2e-16 ***
skew    1.055e+00   2.446e-02   43.141 < 2e-16 ***
shape   7.193e+00   8.408e-01    8.556 < 2e-16 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Log Likelihood:

```
11806.58    normalized:  3.291491
```

Description:

```
Sun Jun 15 09:52:19 2025 by user:
```

Standardised Residuals Tests:

			Statistic	p-Value
Jarque-Bera Test	R	Chi^2	725.5180693	0.00000000
Shapiro-Wilk Test	R	W	0.9831582	0.00000000
Ljung-Box Test	R	Q(10)	10.4216838	0.40430817
Ljung-Box Test	R	Q(15)	27.7618601	0.02310814
Ljung-Box Test	R	Q(20)	29.1478746	0.08488825
Ljung-Box Test	R^2	Q(10)	8.0543302	0.62352994
Ljung-Box Test	R^2	Q(15)	11.0084514	0.75199504
Ljung-Box Test	R^2	Q(20)	15.6980861	0.73516897
LM Arch Test	R	TR^2	10.2583123	0.59331040

Information Criterion Statistics:

AIC	BIC	SIC	HQIC
-6.579637	-6.569291	-6.579642	-6.575949

Title:

```
GARCH Modelling
```

Call:

```
garchFit(formula = ~garch(1, 1), data = res_cac40, cond.dist = "sstd",
  trace = FALSE)
```

Mean and Variance Equation:

```
data ~ garch(1, 1)
<environment: 0x13266e800>
[data = res_cac40]
```

Conditional Distribution:

```
sstd
```

Coefficient(s):

mu	omega	alpha1	beta1	skew	shape
-4.1284e-04	2.0402e-06	6.5401e-02	9.2730e-01	1.0840e+00	1.0000e+01

Std. Errors:

```
based on Hessian
```

Error Analysis:

	Estimate	Std. Error	t value	Pr(> t)
mu	-4.128e-04	2.377e-04	-1.737	0.08245 .
omega	2.040e-06	7.576e-07	2.693	0.00708 **
alpha1	6.540e-02	9.543e-03	6.853	7.23e-12 ***
beta1	9.273e-01	1.030e-02	90.004	< 2e-16 ***
skew	1.084e+00	3.327e-02	32.581	< 2e-16 ***
shape	1.000e+01	1.342e+00	7.452	9.24e-14 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Log Likelihood:

7500.141 normalized: 2.911545

Description:

Sun Jun 15 09:52:20 2025 by user:



Standardised Residuals Tests:

			Statistic	p-Value
Jarque-Bera Test	R	Chi ²	36.6432206	1.104144e-08
Shapiro-Wilk Test	R	W	0.9967943	2.934125e-05
Ljung-Box Test	R	Q(10)	12.0943005	2.787949e-01
Ljung-Box Test	R	Q(15)	20.7596486	1.446860e-01
Ljung-Box Test	R	Q(20)	22.3409037	3.223024e-01
Ljung-Box Test	R ²	Q(10)	12.3893095	2.598459e-01
Ljung-Box Test	R ²	Q(15)	12.9276394	6.078871e-01
Ljung-Box Test	R ²	Q(20)	14.5756952	8.001500e-01
LM Arch Test	R	TR ²	13.4251438	3.389111e-01

Information Criterion Statistics:

AIC	BIC	SIC	HQIC
-5.818432	-5.804797	-5.818443	-5.813490

F Residuals: SP500 ARMA & F Residuals: CAC40 ARMA & (



Box-Ljung test

```
data: residuals(garch_sp500_arma_res)
X-squared = 30.657, df = 20, p-value = 0.05989
```

Box-Ljung test

```
data: residuals(garch_cac40_arma_res)
X-squared = 41.079, df = 20, p-value = 0.003639
```

For SP500, we observe that the ARMA + GARCH fit improves the model quality quite clearly: higher log-likelihood, lower AIC and no autocorrelation

On the contrary, CAC40 has an identical log-likelihood and AIC as before, with a p-value still too low (so still some significant autocorrelation). It could be good to tune ARMA better or to try another GARCH variant.

Question h)

Use the garchAuto.R script in order to fit a GARCH on the residuals of the ARMA(p,q) from (g).
Assess the quality of the fit.

Loading required package: parallel

```
Analyzing (0,0,1,1) with sged distribution done.Good model. AIC = -6.577199, forecast: 0
Analyzing (0,1,1,1) with sged distribution done.Good model. AIC = -6.576938, forecast: 2e-04
Analyzing (0,2,1,1) with sged distribution done.Good model. AIC = -6.576852, forecast: 2e-04
Analyzing (0,3,1,1) with sged distribution done.Good model. AIC = -6.577534, forecast: 1e-04
Analyzing (0,4,1,1) with sged distribution done.Good model. AIC = -6.577077, forecast: 1e-04
Analyzing (0,5,1,1) with sged distribution done.Good model. AIC = -6.578182, forecast: -2e-04
Analyzing (1,0,1,1) with sged distribution done.Good model. AIC = -6.576932, forecast: 2e-04
```


Analyzing (1,1,1,1) with sged distribution done.Bad model.
 Analyzing (1,2,1,1) with sged distribution done.Bad model.
 Analyzing (1,3,1,1) with sged distribution done.Good model. AIC = -6.577338, forecast: 2e-04
 Analyzing (1,4,1,1) with sged distribution done.Good model. AIC = -6.576873, forecast: 2e-04
 Analyzing (1,5,1,1) with sged distribution done.Good model. AIC = -6.578753, forecast: -1e-04
 Analyzing (2,0,1,1) with sged distribution done.Good model. AIC = -6.576831, forecast: 2e-04
 Analyzing (2,1,1,1) with sged distribution done.Bad model.
 Analyzing (2,2,1,1) with sged distribution done.Bad model.
 Analyzing (2,3,1,1) with sged distribution done.Good model. AIC = -6.578384, forecast: -4e-04
 Analyzing (2,4,1,1) with sged distribution done.Bad model.
 Analyzing (2,5,1,1) with sged distribution done.Bad model.
 Analyzing (3,0,1,1) with sged distribution done.Good model. AIC = -6.577463, forecast: 1e-04
 Analyzing (3,1,1,1) with sged distribution done.Good model. AIC = -6.577279, forecast: 2e-04
 Analyzing (3,2,1,1) with sged distribution done.Good model. AIC = -6.578381, forecast: -4e-04
 Analyzing (3,3,1,1) with sged distribution done.Bad model.
 Analyzing (3,4,1,1) with sged distribution done.Bad model.
 Analyzing (3,5,1,1) with sged distribution done.Bad model.
 Analyzing (4,0,1,1) with sged distribution done.Good model. AIC = -6.577022, forecast: 2e-04
 Analyzing (4,1,1,1) with sged distribution done.Good model. AIC = -6.576817, forecast: 2e-04
 Analyzing (4,2,1,1) with sged distribution done.Bad model.
 Analyzing (4,3,1,1) with sged distribution done.Bad model.
 Analyzing (4,4,1,1) with sged distribution done.Bad model.
 Analyzing (4,5,1,1) with sged distribution done.Bad model.
 Analyzing (5,0,1,1) with sged distribution done.Good model. AIC = -6.578058, forecast: -2e-04
 Analyzing (5,1,1,1) with sged distribution done.Good model. AIC = -6.578805, forecast: -2e-04
 Analyzing (5,2,1,1) with sged distribution done.Bad model.
 Analyzing (5,3,1,1) with sged distribution done.Bad model.
 Analyzing (5,4,1,1) with sged distribution done.Bad model.
 Analyzing (5,5,1,1) with sged distribution done.Bad model.

Analyzing (0,0,1,1) with sged distribution done.Good model. AIC = -5.820957, forecast: -4e-04
 Analyzing (0,1,1,1) with sged distribution done.Good model. AIC = -5.820384, forecast: -4e-04
 Analyzing (0,2,1,1) with sged distribution done.Good model. AIC = -5.820133, forecast: -5e-04
 Analyzing (0,3,1,1) with sged distribution done.Good model. AIC = -5.822429, forecast: -6e-04
 Analyzing (0,4,1,1) with sged distribution done.Good model. AIC = -5.82211, forecast: -7e-04
 Analyzing (0,5,1,1) with sged distribution done.Good model. AIC = -5.824134, forecast: -4e-04
 Analyzing (1,0,1,1) with sged distribution done.Good model. AIC = -5.820385, forecast: -4e-04
 Analyzing (1,1,1,1) with sged distribution done.Bad model.
 Analyzing (1,2,1,1) with sged distribution done.Good model. AIC = -5.821597, forecast: -0.0014
 Analyzing (1,3,1,1) with sged distribution done.Bad model.
 Analyzing (1,4,1,1) with sged distribution done.Bad model.
 Analyzing (1,5,1,1) with sged distribution done.Good model. AIC = -5.824193, forecast: -8e-04
 Analyzing (2,0,1,1) with sged distribution done.Good model. AIC = -5.820152, forecast: -5e-04
 Analyzing (2,1,1,1) with sged distribution done.Good model. AIC = -5.821675, forecast: -0.0015
 Analyzing (2,2,1,1) with sged distribution done.Bad model.
 Analyzing (2,3,1,1) with sged distribution done.Good model. AIC = -5.823722, forecast: -7e-04
 Analyzing (2,4,1,1) with sged distribution done.Good model. AIC = -5.823997, forecast: -0.001
 Analyzing (2,5,1,1) with sged distribution done.Good model. AIC = -5.824055, forecast: -2e-04
 Analyzing (3,0,1,1) with sged distribution done.Good model. AIC = -5.822217, forecast: -5e-04
 Analyzing (3,1,1,1) with sged distribution done.Good model. AIC = -5.821607, forecast: 4e-04
 Analyzing (3,2,1,1) with sged distribution done.Good model. AIC = -5.823868, forecast: -8e-04
 Analyzing (3,3,1,1) with sged distribution done.Bad model.

```

Analyzing (3,4,1,1) with sged distribution done.Bad model.
Analyzing (3,5,1,1) with sged distribution done.Bad model.
Analyzing (4,0,1,1) with sged distribution done.Good model. AIC = -5.821934, forecast: -5e-04
Analyzing (4,1,1,1) with sged distribution done.Good model. AIC = -5.822103, forecast: 3e-04
Analyzing (4,2,1,1) with sged distribution done.Bad model.
Analyzing (4,3,1,1) with sged distribution done.Bad model.
Analyzing (4,4,1,1) with sged distribution done.Bad model.
Analyzing (4,5,1,1) with sged distribution done.Bad model.
Analyzing (5,0,1,1) with sged distribution done.Good model. AIC = -5.823899, forecast: -2e-04
Analyzing (5,1,1,1) with sged distribution done.Good model. AIC = -5.824684, forecast: -7e-04
Analyzing (5,2,1,1) with sged distribution done.Bad model.
Analyzing (5,3,1,1) with sged distribution done.Bad model.
Analyzing (5,4,1,1) with sged distribution done.Bad model.
Analyzing (5,5,1,1) with sged distribution done.Bad model.

```

Title:

GARCH Modelling

Call:

```

garchFit(formula = formula, data = data, cond.dist = ll$dist,
  trace = FALSE)

```

Mean and Variance Equation:

data ~ arma(5, 1) + garch(1, 1)

<environment: 0x125f760c0>

[data = data]

Conditional Distribution:

sged

Coefficient(s):

mu	ar1	ar2	ar3	ar4	ar5
-2.2520e-05	6.1560e-01	-5.1324e-05	-1.9771e-02	2.1914e-02	-3.3774e-02
ma1	omega	alpha1	beta1	skew	shape
-6.3161e-01	3.8994e-07	5.2238e-02	9.4520e-01	1.0729e+00	1.3954e+00

Std. Errors:

based on Hessian

Error Analysis:

	Estimate	Std. Error	t value	Pr(> t)
mu	-2.252e-05	5.003e-05	-0.450	0.652641
ar1	6.156e-01	1.885e-01	3.265	0.001093 **
ar2	-5.132e-05	2.043e-02	-0.003	0.997996
ar3	-1.977e-02	1.974e-02	-1.001	0.316667
ar4	2.191e-02	2.052e-02	1.068	0.285454
ar5	-3.377e-02	1.918e-02	-1.761	0.078316 .
ma1	-6.316e-01	1.888e-01	-3.345	0.000823 ***
omega	3.899e-07	1.556e-07	2.506	0.012194 *
alpha1	5.224e-02	7.840e-03	6.663	2.69e-11 ***
beta1	9.452e-01	7.986e-03	118.360	< 2e-16 ***
skew	1.073e+00	2.364e-02	45.383	< 2e-16 ***

shape 1.395e+00 4.584e-02 30.442 < 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Log Likelihood:

11811.09 normalized: 3.292748

Description:

Sun Jun 15 09:53:20 2025 by user:

Standardised Residuals Tests:

			Statistic	p-Value
Jarque-Bera Test	R	Chi ²	766.0267687	0.000000000
Shapiro-Wilk Test	R	W	0.9822839	0.000000000
Ljung-Box Test	R	Q(10)	14.5209369	0.150528097
Ljung-Box Test	R	Q(15)	31.5502756	0.007408813
Ljung-Box Test	R	Q(20)	32.8242579	0.035269235
Ljung-Box Test	R ²	Q(10)	7.8181492	0.646594554
Ljung-Box Test	R ²	Q(15)	11.0400638	0.749749175
Ljung-Box Test	R ²	Q(20)	15.6369602	0.738875811
LM Arch Test	R	TR ²	10.1047177	0.606774224

Information Criterion Statistics:

AIC	BIC	SIC	HQIC
-6.578805	-6.558113	-6.578827	-6.571430

Title:

GARCH Modelling

Call:

```
garchFit(formula = formula, data = data, cond.dist = ll$dist,
trace = FALSE)
```

Mean and Variance Equation:

```
data ~ arma(5, 1) + garch(1, 1)
<environment: 0x135cf4698>
[data = data]
```

Conditional Distribution:

sged

Coefficient(s):

mu	ar1	ar2	ar3	ar4	ar5
-1.8718e-04	6.2888e-01	-1.3514e-02	-4.6209e-02	3.3830e-02	-5.3382e-02
ma1	omega	alpha1	beta1	skew	shape
-6.3026e-01	2.0199e-06	6.5264e-02	9.2512e-01	1.1129e+00	1.7928e+00

Std. Errors:

based on Hessian

Error Analysis:

	Estimate	Std. Error	t value	Pr(> t)
mu	-1.872e-04	1.063e-04	-1.761	0.07831 .
ar1	6.289e-01	1.329e-01	4.731	2.23e-06 ***
ar2	-1.351e-02	2.371e-02	-0.570	0.56869
ar3	-4.621e-02	2.397e-02	-1.928	0.05390 .
ar4	3.383e-02	2.492e-02	1.357	0.17469
ar5	-5.338e-02	2.107e-02	-2.534	0.01127 *
ma1	-6.303e-01	1.317e-01	-4.784	1.72e-06 ***
omega	2.020e-06	7.125e-07	2.835	0.00458 **
alpha1	6.526e-02	8.804e-03	7.413	1.23e-13 ***
beta1	9.251e-01	9.948e-03	92.996	< 2e-16 ***
skew	1.113e+00	3.255e-02	34.186	< 2e-16 ***
shape	1.793e+00	7.472e-02	23.993	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Log Likelihood:
7514.193 normalized: 2.917

Description:
Sun Jun 15 09:54:04 2025 by user:


Standardised Residuals Tests:

			Statistic	p-Value
Jarque-Bera Test	R	Chi^2	42.7453156	5.223633e-10
Shapiro-Wilk Test	R	W	0.9961469	3.556087e-06
Ljung-Box Test	R	Q(10)	8.3033881	5.992288e-01
Ljung-Box Test	R	Q(15)	17.6818769	2.797593e-01
Ljung-Box Test	R	Q(20)	19.6727478	4.785610e-01
Ljung-Box Test	R^2	Q(10)	13.5760142	1.932238e-01
Ljung-Box Test	R^2	Q(15)	14.0134143	5.245120e-01
Ljung-Box Test	R^2	Q(20)	16.4207731	6.902011e-01
LM Arch Test	R	TR^2	14.2477321	2.851692e-01

Information Criterion Statistics:

AIC	BIC	SIC	HQIC
-5.824684	-5.797414	-5.824727	-5.814799

ACF Residuals: SP500 Auto GARCH Residuals: CAC40 Auto G



Box-Ljung test

```
data: residuals(best_garch_sp500)
X-squared = 28.262, df = 20, p-value = 0.1033
```

Box-Ljung test

```
data: residuals(best_garch_cac40)
X-squared = 22.765, df = 20, p-value = 0.3005
```

As a result, both SP500 and CAC40 now have the best fit compared to the previous fits we have tried. Residuals are now white noise for both indices (so no more autocorrelation) and volatility are well-captured.


Practical 2, Part 1 - Venice

The `venice90` dataset can be found in the `VGAM` package.

Question a)

Read in the data. Extract and represent the yearly max values from 1940 to 2009. What do you observe ?

Yearly Maximum Sea Levels in Venice (1940–2009)




We can observe some variability over the years and a slight upward trend, so the maximum levels in Venice seem to be increasing.

Question b)

We are end of 2009 and would like to predict the yearly maximum values over the next 13 years (from 2010 to 2022). A naive approach consists of fitting a linear model on the observed yearly maxima and predict their values for 2010–2022. Proceed to this prediction and provide confidence intervals.

	year	fit	lwr	upr
1	2010	132.4522	121.4683	143.4361
2	2011	132.7321	121.5137	143.9505
3	2012	133.0121	121.5576	144.4665
4	2013	133.2920	121.6002	144.9838
5	2014	133.5719	121.6414	145.5025
6	2015	133.8519	121.6813	146.0225
7	2016	134.1318	121.7200	146.5436
8	2017	134.4118	121.7576	147.0659
9	2018	134.6917	121.7942	147.5892
10	2019	134.9716	121.8298	148.1135
11	2020	135.2516	121.8644	148.6387
12	2021	135.5315	121.8982	149.1649
13	2022	135.8115	121.9311	149.6919

Observed and Predicted Yearly Maximum Sea Levels




We used a confidence interval of 99% to predict for the years 2010 to 2022.

Question c)

Represent in the same graph the predicted yearly max values for the period 2010–2022, their point-wise confidence bounds and the observed values greater than 140 cm from the table below.

Predicted Max Sea Levels (2010–2022) with Historical Extremes



This plot provides all the necessary information, from the historical data in the blue line, to the yearly maximum values with the red points, the dark green line being the prediction for 2010 to 2022, the light green area being the confidence intervals and finally, the black triangles being the values greater than 140cm.

Now we perform a risk analysis and because we are interested in the period 2010–2022, we want to calculate the 13-years return level., for each year.

Question d)


Fit a GEV a with constant parameters to the historical yearly max values. Fit a GEV with time varying location parameter. Compare the two embedded models using likelihood ratio test (LRT). Show diagnostic plots.

--- d) Likelihood-ratio test ---

LRT = 11.62, $p = 0.001$

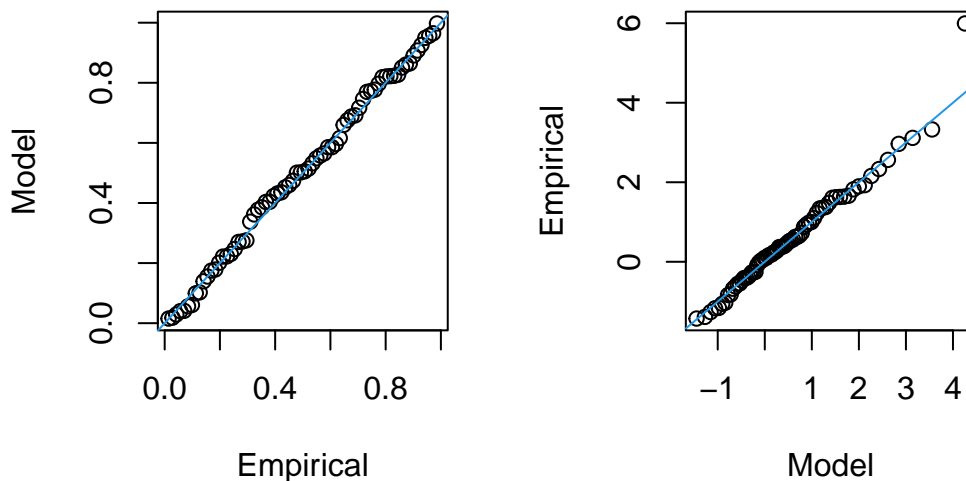
Selected model: Time-varying Location

Constant Parameters



Time-varying Location

Residual Probability Plot & Residual Quantile Plot (Gumbel)



We fitted a constant and a time-varying model. The latter is better thanks to the low p-value and the log-likelihood of 11.62. The model looks overall okay, despite having some outliers which might influence it. There are no major pattern nor heteroskedasticity.

Question e)

Add if necessary a time varying scale and or shape GEV parameter. Select the best model according to LRT.

```
# A tibble: 4 x 4
  comparison      LR      df      p
  <chr>      <dbl> <dbl> <dbl>
1 location vs const 11.6      1 0.000651
2 location+scale vs location 0.892      1 0.345
3 location+shape vs location 5.03      1 0.0250
4 location+scale+shape vs location+scale 5.94      1 0.0148
```

Selected model: location+shape

The best model includes time-varying location and shape parameters. The addition of a time-varying scale is not necessary based on the LRT. This model provides the best fit and should be used for further analysis or prediction.

Question f) + g)

- f) Predict the 13-years return level, each year from 2010 to 2022.
- g) Calculate confidence bands for these predictions.

	year	return_level	lower_bound	upper_bound
1	2010	147.78	137.25	158.32
2	2011	147.87	137.72	158.02
3	2012	147.96	138.17	157.75
4	2013	148.06	138.60	157.52
5	2014	148.17	139.00	157.33
6	2015	148.27	139.38	157.17
7	2016	148.39	139.73	157.04
8	2017	148.51	140.07	156.94
9	2018	148.63	140.39	156.87
10	2019	148.75	140.68	156.82
11	2020	148.88	140.96	156.81
12	2021	149.02	141.22	156.82
13	2022	149.16	141.47	156.85


For each year from 2010 to 2022, the estimated 13-year return level gradually increases from approximately 147.78 cm to 149.16 cm. This indicates a slight upward trend in extreme sea level risk over time. The 95% confidence intervals range from about 137–158 cm in 2010 to 141–157 cm in 2022, showing that while uncertainty remains, the expected extremes are becoming higher. This trend supports the idea that extreme sea level events in Venice are becoming more likely and potentially more severe over time.

Question h)

Represent in the same graph your predictions of the 13-years return levels, their pointwise confidence intervals, the predicted yearly max values from the linear model and the observed values greater than 140 cm from the table below.

Venice Yearly Maxima, Forecasts, and 13–Year Return Levels

Blue = Observed (1940–2009) · Green = Linear Model Forecast · Red = 1



Question i)

Broadly speaking, each year, there is a chance of $1/13$ that the observed value is above the 13-years return level. Comment the results for both the linear model prediction and GEV approach. Note that 12 of the 20 events occurred in the 21st century.


While both models provide useful insights, the linear model clearly underestimates extremes and provides overly narrow confidence intervals. The GEV approach, especially with time-varying parameters, is more suited for modeling extremes and gives a more realistic picture of sea level risk. However, even the GEV predictions fall short of the most recent high events, such as 2.04m in 2022, indicating that the system is non-stationary and that risk is increasing over time. This shift is emphasized by the concentration of extreme events in the 21st century, suggesting that return periods are shortening and that what was once a 13-year event may now be happening more frequently.

Practical 2, Part 2 - Nuclear Reactors

Question a)

Read in the data. Display a time series plot of the water level across the data range and try to identify times of highest levels.

Daily Maximum Water Level Over Time

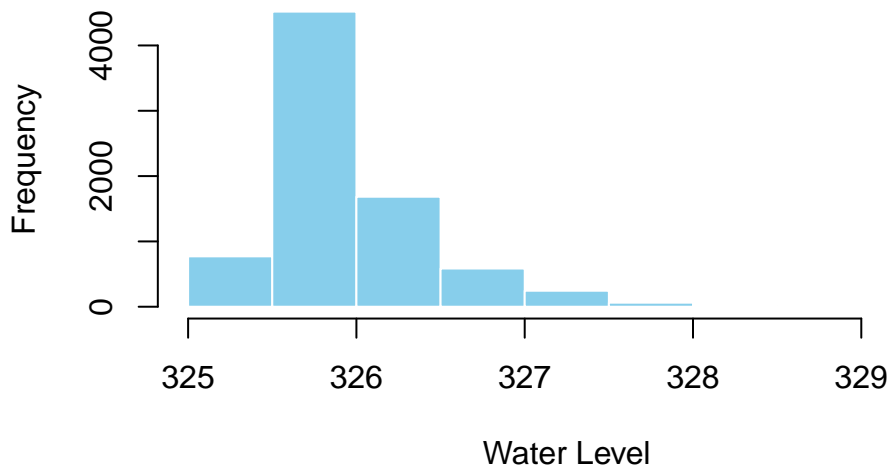


	Stationsname	Stationsnummer	Parameter	Zeitreihe	
2778	Untersiggenthal, Stilli	2205	Pegel	Tagesmaxima	
2061	Untersiggenthal, Stilli	2205	Pegel	Tagesmaxima	
2062	Untersiggenthal, Stilli	2205	Pegel	Tagesmaxima	
7865	Untersiggenthal, Stilli	2205	Pegel	Tagesmaxima	
7866	Untersiggenthal, Stilli	2205	Pegel	Tagesmaxima	
	Parametereinheit	Gewässer	Zeitstempel	Zeitpunkt_des_Auftretens	Wert
2778	m ü.M.	Aare	2007-08-09	2007-08-09 11:15:00	329.323
2061	m ü.M.	Aare	2005-08-22	2005-08-22 18:05:00	328.827
2062	m ü.M.	Aare	2005-08-23	2005-08-23 08:35:00	328.742
7865	m ü.M.	Aare	2021-07-14	2021-07-14 07:45:00	328.622
7866	m ü.M.	Aare	2021-07-15	2021-07-15 15:05:00	328.614
	Freigabestatus				
2778	Freigegeben, validierte Daten				
2061	Freigegeben, validierte Daten				
2062	Freigegeben, validierte Daten				
7865	Freigegeben, provisorische Daten				
7866	Freigegeben, provisorische Daten				

Question b)

Now display a histogram of the water levels. What do you observe about the distribution?

Histogram of Water Levels



The distribution is right-skewed. Most levels are concentrated between 325 and 326. Extreme levels such as above 327 are rare yet still present. These can represent potential flood events or unusual conditions.

The FOEN plans for several degrees of risk. In this assignment, we focus on two risk levels: 50-year events and 100-year events.

Question c)

Explain how you would model the high water levels using a peaks-over-threshold approach.

99% threshold: 327.5054

Number of exceedances: 79

Using a Peaks-over-Threshold approach, we set a threshold above which the values are considered extreme. This threshold should be high enough to focus only on rare exceedances, but not too high to avoid having too few exceedances. Here, the threshold is set at the 99th percentile, which is at 327.5054 (so around 327.51) meters. The exceedances are modeled using the Generalized Pareto Distribution, suitable for a skewed distribution. In this case, the POT approach is useful as there are a lot of non-extreme values. This approach thus focuses only on the extreme events to assess a better statistical efficiency, especially with daily data over many years.

Question d)

Comment on the aspect of clustering of extremes. How do you propose to measure and deal with clustering of the daily water levels?

Number of exceedances before declustering: 79

Number of cluster maxima (after declustering): 28

Clustering extremes uses runs methods. We keep only one peak per cluster, which makes the exceedances more independent and suitable for modelling.

Question e)

Perform the analysis you suggest in c) and d) and compute the 50- and 100-year return levels. Explain your choice of threshold and provide an estimate of uncertainty for the return levels. Note: take care to compute the return level in yearly terms.

Using the POT approach:

```
$threshold
```

```
99%
```

```
327.5054
```

```
$nexc
```

```
[1] 28
```

```
$conv
```

```
[1] 0
```

```
$nllh
```

```
[1] 0.8853974
```

```
$mle
```

```
[1] 0.3366714 0.1202787
```

```
$rate
```

```
[1] 1
```

```
$se
```

```
[1] 0.1009116 0.2337588
```

```
50-year return level: 329.33
```

```
95% CI: 328.37 - 331.62
```

```
100-year return level: 329.73
```

```
95% CI: 328.43 - 333.53
```

The threshold is the 99th percentile to capture the extremes, have a balance between bias and variance and to have an adequate sample size to fit a GPD

Question f)

Explain the drawbacks and advantages of using a block maxima method instead of the one used in c)-e).

The Block Maxima method selects the maximum observation from a given time interval, but uses only one observation per block which leads to an inefficient use of the data. The POT approach uses all the values above the given threshold and handles clustering well via declustering. It is more efficient and flexible especially when extreme events happen in clusters. Thus, the POT approach is more precise and provides more information on the behavior of extreme events.

Practical 2, Part 3 - Night temperatures in Lausanne


Question a)

Read in the data for the daily night maximum temperatures in Lausanne. Subset the summer months (June to September).

New names:

New names:

```
* `` -> `...1`
```



We are doing the same process for minimum to answer question e.

Question b)

Assess whether extremes of the subsetting series in (a) occur in cluster.

Runs Estimator for the Extremal Index

extremal.index	number.of.clusters	run.length
0.4019608	41.0000000	3.0000000


Runs Estimator for the Extremal Index

extremal.index	number.of.clusters	run.length
0.3673469	36.0000000	3.0000000

The obtained extremal index is 0.402, which is lower than 1. This suggests that extreme night temperatures during summer in Lausanne tend to occur in clusters rather than being isolated. This means that if you observe one extremely hot night, there is a higher chance that other extreme nights will follow shortly, such as during a heatwave for example.

Question c)

Deccluster the data from (a) using a suitable threshold. Plot the resulting declustered data. (Hint: you may want to use the `extRemes` package.)




After declustering the extreme summer night temperatures using the 95th percentile threshold and a 3-day run length, we isolated 42 independent exceedances above the threshold. The resulting plot of declustered excesses reveals a wide range of magnitudes, with some cluster peaks exceeding 4°C above the threshold. This confirms the presence of significant and varied extreme temperature events, now stripped of temporal dependence.


Question d)

Fit a GPD to the data, both raw and declustered. Assess the quality of the fit.

-Plot: Raw Summer Maxima (left): Declustered Summer Maxima (right)



Q-Plot: Raw Winter Minima (Plot: Declustered Winter Minima)



Despite the raw model having points more aligned in the QQ-plot, the declustered model is theoretically better due to the lower AIC. The deviations in the QQ-plot for the declustered can be explained due to the smaller sample than the raw data.

Question e)

Repeat the above analysis for the negatives of the daily nightly minimum temperatures for the winter months (November-February).

```
$winter_extremal_index
```

```
Runs Estimator for the Extremal Index
```

extremal.index	number.of.clusters	run.length
0.3673469	36.0000000	3.0000000

```
$gpd_winter_raw
```

```
fevd(x = (-winter_min$tmin)[-winter_min$tmin > u_min] - u_min,
      threshold = 0, type = "GP", method = "MLE")
```

```
[1] "Estimation Method used: MLE"
```

```
Negative Log-Likelihood Value: 186.0393
```

```
Estimated parameters:
```

scale	shape
2.8113093	-0.1352909

```
Standard Error Estimates:
```

scale	shape

0.4203671 0.1110884

Estimated parameter covariance matrix.

	scale	shape
scale	0.17670849	-0.03804989
shape	-0.03804989	0.01234063

AIC = 376.0786

BIC = 381.2486

\$gpd_winter_dc

fevd(x = excess_min, threshold = 0, type = "GP", method = "MLE")

[1] "Estimation Method used: MLE"

Negative Log-Likelihood Value: 68.36057

Estimated parameters:

	scale	shape
	3.0113921	-0.2034978

Standard Error Estimates:

	scale	shape
	0.6471591	0.1398393


Estimated parameter covariance matrix.

	scale	shape
scale	0.41881484	-0.07244737
shape	-0.07244737	0.01955502

AIC = 140.7211

BIC = 143.8882


Summer Daily Minimum Water Levels (November–February, 2000–2022)



Declassified Winter Night Temperature Excesses



Q-Plot: Raw Winter Minima (Plot: Declustered Winter Minima)



We apply the negative to the winter values to treat the extremely low values as high values for modelling purposes. We then do an extremal index and we obtain 0.367, lower than 1, indicating clustering. We then declustered using the 95th percentile threshold to the negated temperatures and the plot shows that some peaks go even above 6 degrees. Fitting the model using GPD shows that the AIC for the declustered is again lower than raw, so a better fit.